FIDE AI / AGENTDOJO AUDIT

Same reconstructed states. Different scores.

Direct task-0 predicate substitution on the same 29 tasks in each configuration.

The ordering stays the same. These are reconstructed scores, not fresh trials.

Later event-evidence analysis leaves repeat prompt at 5–6/29; point five retains the original state-reconstruction assumptions. Event bounds are a different measure, not confidence intervals.

Source: pinned AgentDojo records; Fide predicate comparison. All values are counts out of 29.