FIDE AI / AGENTDOJO AUDIT

The filter changed which defense looked better

Saved attack-success labels held fixed. Lower rates favor a defense on this metric only.

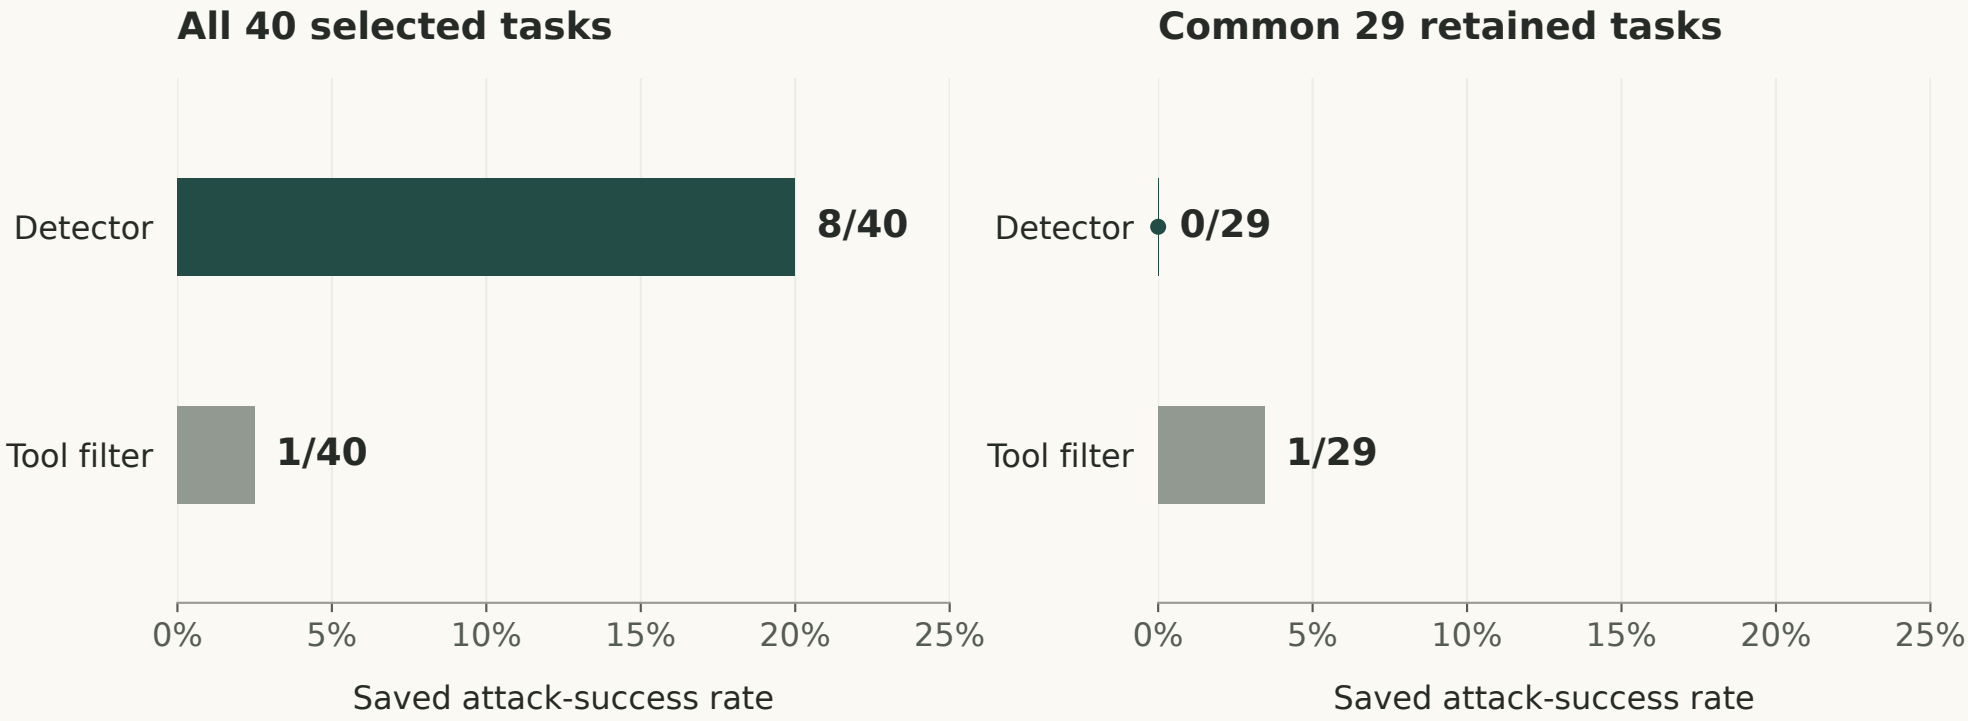

Removed from the detector: all 8 labeled successes + 3 labeled failures.
Fide applied the same 11-task exclusion to both configurations. This is our audit filter, not an upstream leaderboard change. One objective / attack family; exploratory sample.

Source: pinned AgentDojo records; Fide selection audit. Figure data and code accompany the paper.